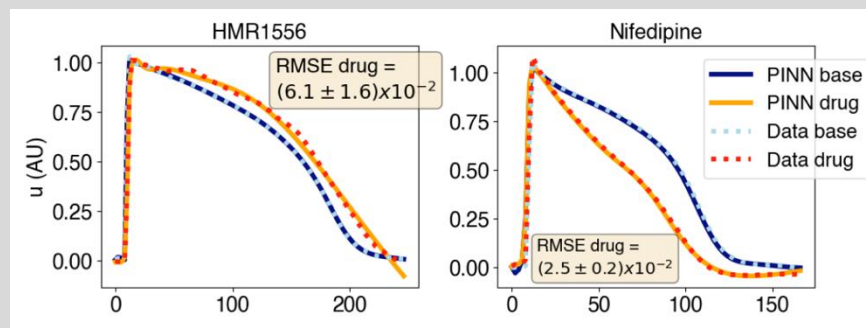
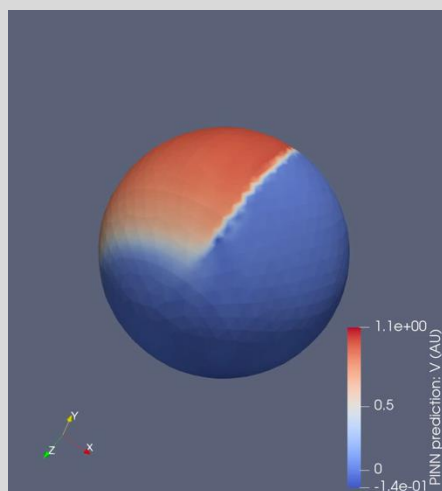


# PINNs Tutorial

- Introduction to PINNs
  - What are PINNs?
  - Considerations when designing PINNs
- Google Colab Exercises
  - PINNs to estimate cerebral blood flow
  - PINNs to model cardiac electrophysiology
- Recent developments in PINNs
  - Problems with PINNs
  - How to address them



# Introduction to Physics-Informed Neural Networks



**Annie Cing-En Chiu & Marta Varela**

[ching-en.chiu18@imperial.ac.uk](mailto:ching-en.chiu18@imperial.ac.uk)

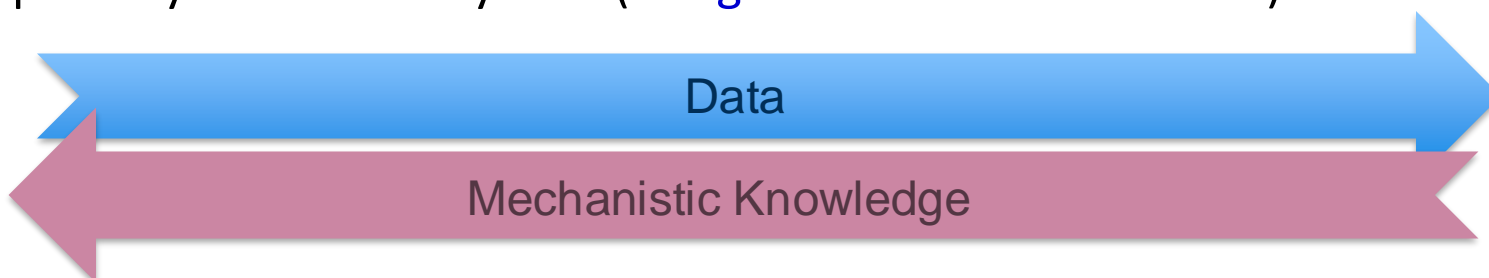
[marta.varela@imperial.ac.uk](mailto:marta.varela@imperial.ac.uk)

**IMPERIAL**

# Physics-Informed Machine Learning

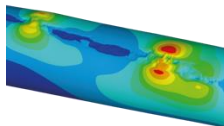
## Aims

- Model and understand the physiology of the system (**organ, tissue**)
- Identify, predict and characterise failure modes (**pathology**)
- Optimally control the system (**design and assess treatments**)



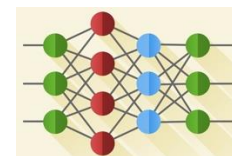
- Differential equations known
- All parameters identified
- Initial and boundary conditions known

- A lot of data is available
- Data perfectly reproduces the system and is representative of the population



Finite elements,  
finite differences ...

Neural networks ...

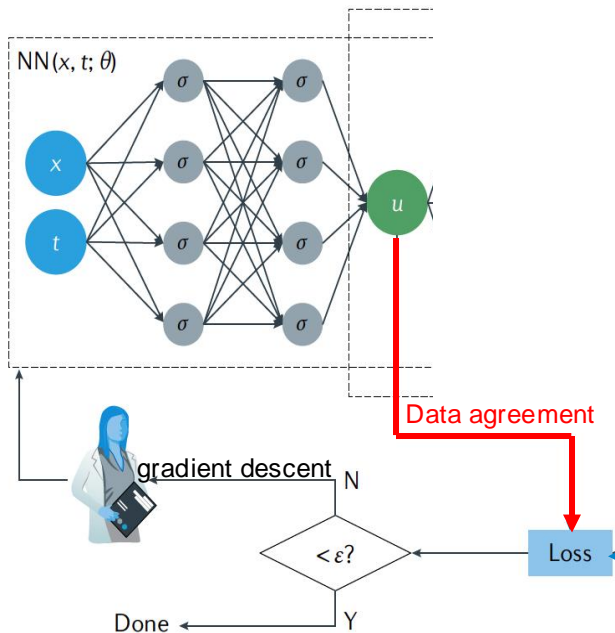


Physics-Informed Machine  
Learning (e.g., PINNs)

# Physics-Informed Neural Networks (PINNs)



$$\frac{\partial u}{\partial t} = D \frac{\partial^2 u}{\partial x^2} \quad \text{Diffusion Equation}$$



$$\mathcal{L} = w_{\text{data}} \mathcal{L}_{\text{data}} + w_{\text{PDE}} \mathcal{L}_{\text{PDE}},$$

where

$$\mathcal{L}_{\text{data}} = \frac{1}{N_{\text{data}}} \sum_{i=1}^{N_{\text{data}}} (u(x_i, t_i) - u_i)^2 \quad \text{and}$$

$$\mathcal{L}_{\text{PDE}} = \frac{1}{N_{\text{PDE}}} \sum_{j=1}^{N_{\text{PDE}}} \left( \frac{\partial u}{\partial t} - D \frac{\partial^2 u}{\partial x^2} \right)^2 \Big|_{(x_j, t_j)}.$$

## Forward Mode

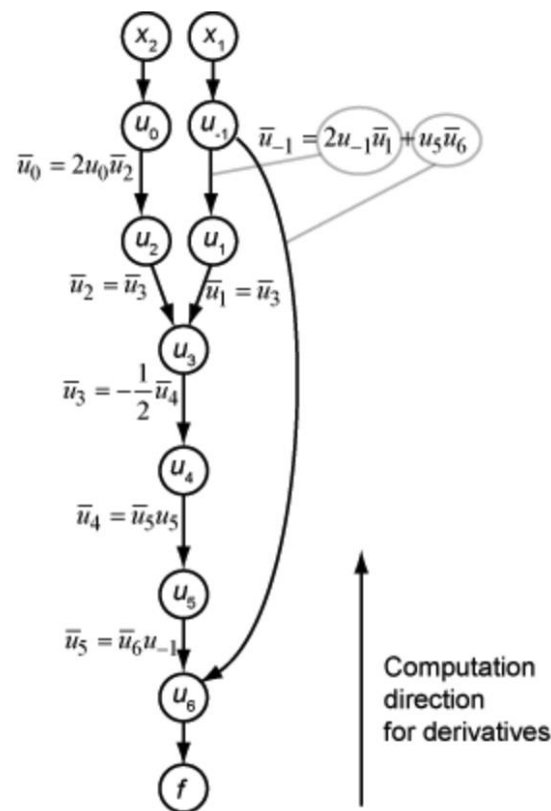
- NN to solve the differential equation in entire domain (find function  $u(x, t)$ )
- No discretization of domain performed

## Inverse Mode

- Also estimate equation parameters, e.g. diffusion coefficient,  $D$
- Adjusted using gradient descent

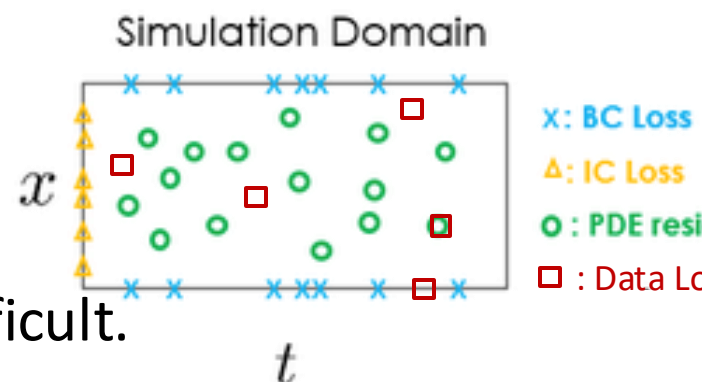
# Calculate Derivatives Fast: Auto Diff

- Automatic Differentiation (**AutoDiff**):
  - The workhorse of backpropagation for NN training
  - More widely applicable than **symbolic diff**
  - No discretization errors, discrete solutions as in **numerical diff**
  - Faster, especially for high-order and multivariate diff
- Based on the algebra of dual numbers and the chain rule.
- In the PINNs framework, equation derivatives come for free!



# Count Your Losses

- Loss terms typically use the mean squared error (MSE).
- Agreement with equations is computed at pre-assigned points ( $N_{\text{PDE/ODE}}$ ). Unrelated to the number of data points ( $N_{\text{data}}$ ).
- Different loss terms can have different weights.
- Can also adaptatively refine location of points where  $L_{\text{PDE/ODE}}$  is calculated to minimise it<sup>1</sup>.
- PINNs can work in forward mode with no experimental data (just BCs, ICs), but...
  - Convergence is slow(er) and (more) difficult.
  - If you have extra data, use it!



# On the Edge: Initial & Boundary Conditions

- Initial and/or boundary conditions should be included in the loss function (soft enforcement).

$$\mathcal{L} = \mathcal{L}_{PDE/ODE} + \mathcal{L}_{IC} + \mathcal{L}_{BC} + \mathcal{L}_{data}$$

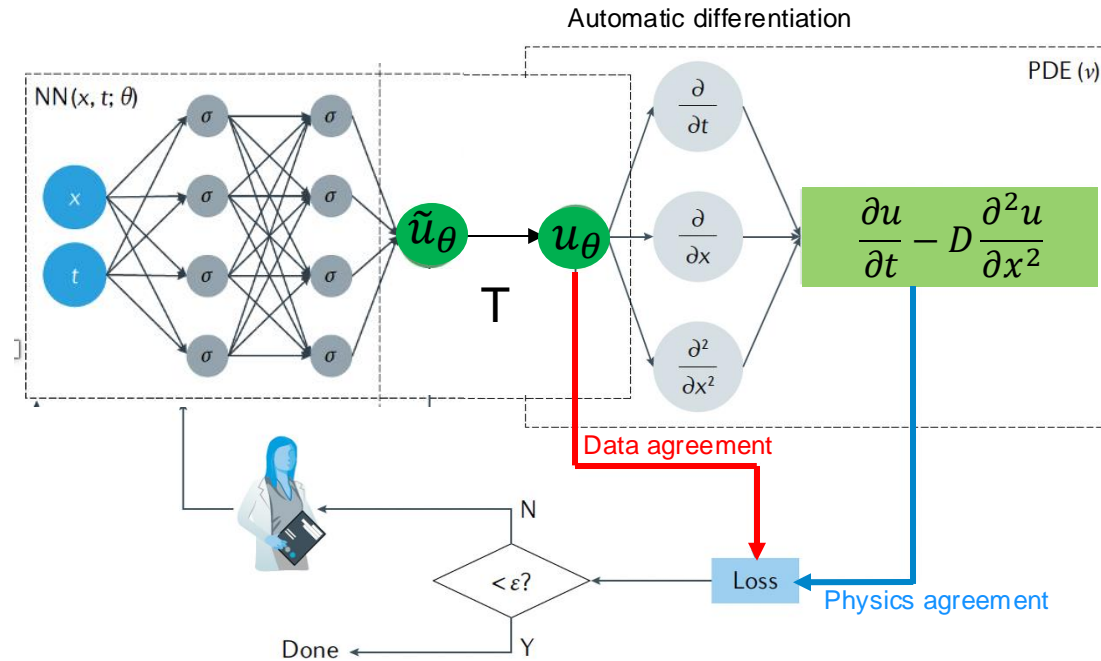
- The IC and BC losses resemble data losses in specific conditions.

$$\mathcal{L}_{IC} = \frac{1}{N_{ptsIC}} \sum_{i=1}^{N_{ptsIC}} (u_{\theta}(x_i, 0) - u(x_i, 0))$$

$$\mathcal{L}_{BC} = \frac{1}{N_{ptsBC}} \sum_{i=1}^{N_{ptsBC}} (u_{\theta}(x_{BC}, t_i) - u(x_{BC}, t_i))$$

- For hard enforcement, we can apply transforms to the NN solutions so that ICs and BCs are obeyed.

# Hard BCs, ICs



- Examples:

$$\frac{\partial y}{\partial t} = \frac{\partial^2 y}{\partial x^2} - e^{-t}(\sin(\pi x) - \pi^2 \sin(\pi x)), \quad x \in [-1, 1], \quad t \in [0, 1]$$

with  $y(x, 0) = \sin(\pi x)$  and  $y(-1, t) = y(1, t) = 0$ .

gives e.g.,  $T = t * (1 - x^2) * y + \sin(\pi x)$

Example taken from:

[https://deepxde.readthedocs.io/en/latest/demos/pinn\\_forward/diffusion.1d.exactBC.html](https://deepxde.readthedocs.io/en/latest/demos/pinn_forward/diffusion.1d.exactBC.html)

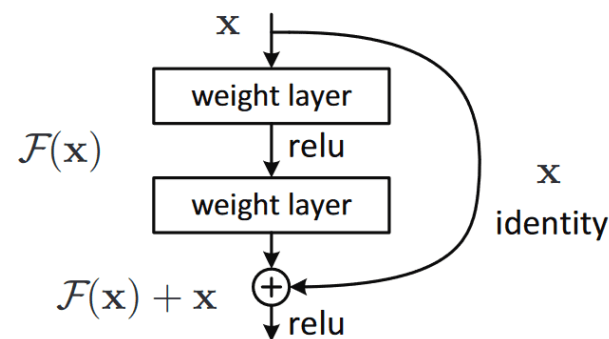
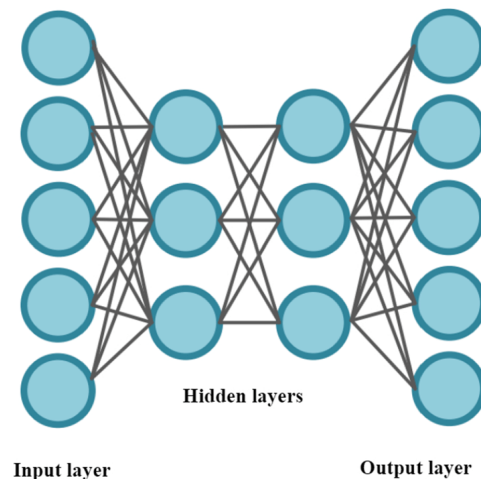


# Build It: Architecture for PINNs

- Fully-connected NNs (= multilayer perceptrons, MLPs) are most popular.

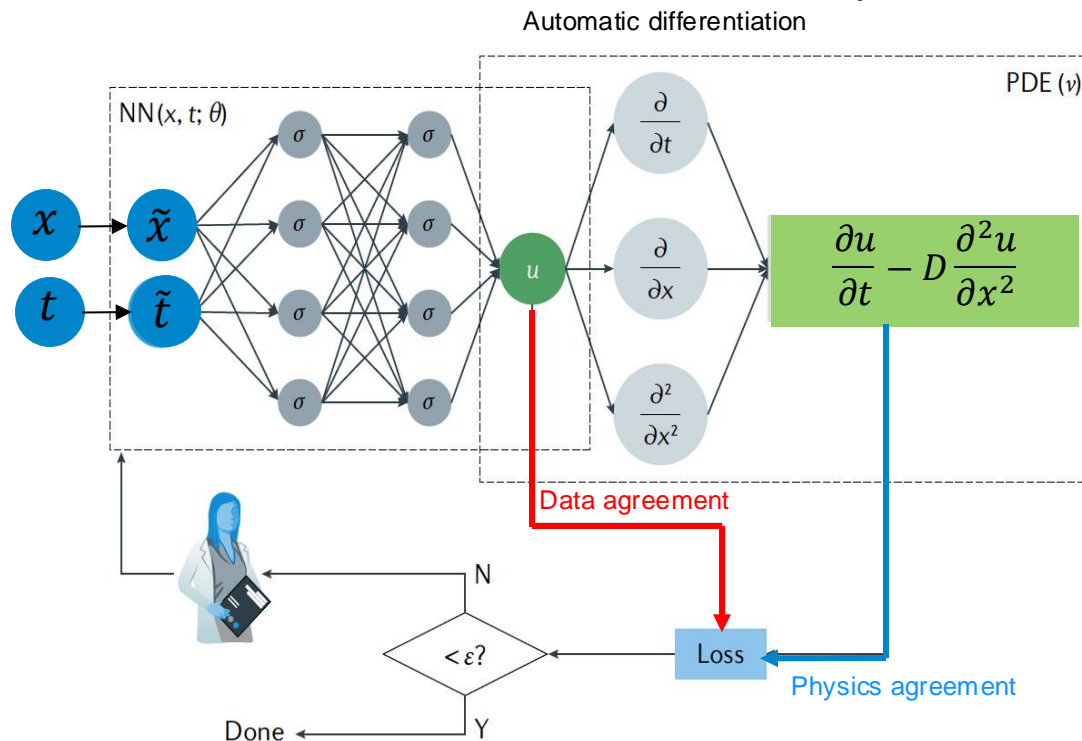
## Problems:

- Vanishing gradients for deep networks
  - ReLU as activation function not commonly used (problems differentiating it!)
  - Tanh (and Swish) most often used
- Skip (=residual) connections sometimes added
- Normalisation of inputs, outputs helps balance gradients
- Depth and width adjusted manually
  - 2-6 layers
  - 32-516 neurons/layer



# High-Frequency Modes

- FCNNs suffer from spectral bias: they preferentially learn low-frequency functions.
  - Can be mitigated by changing activation function
  - In PINNs, more often addressed by applying a (non-learned) Fourier transform to the inputs



## Fourier Mapping

$$\tilde{t} = \sum_{k=1}^{N_{modes}} \sin(kt)$$

## Random Fourier Mapping

$$\tilde{t} = [\sin(Bt), \cos(Bt)]$$

with  $B \in \mathbb{R}^{N_{modes}}$ ,  $B_k \sim \mathcal{N}(0, \sigma^2)$

# Choose the Best: Optimiser for PINNs

- Adam is most commonly used.
  - Low learning rates work best. (Can be reduced iteratively if needed.)
- An L-BFGS pass following Adam sometimes used.
- Batch learning possible in PINNs, but only useful if a lot of training points are used.
  - Full-batch learning is very common.
- NN parameters are typically initialised using Glorot initialisation, although statistical properties of PINNs ‘samples’ are hard to calculate.

# Inverse Mode

- PINNs can also estimate one or more unknown model parameters given some experimental data for the system.
- These parameters are included in the AutoDiff-based optimisation, playing a similar role to network parameters.
- It is a good idea to perform identifiability analyses before running PINNs in inverse mode.

# Make Your Life Easier: Libraries for PINNs

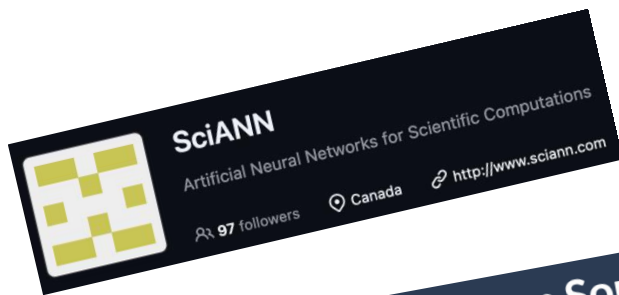
- Here, we will use an open-source Python library (with a TensorFlow backend):
- Other PINN libraries (Python, Julia, Jax):

**DeepXDE 1.12.1**

```
pip install DeepXDE
```



**NVIDIA Modulus**



**SciML: Open Source Software for Scientific Machine Learning**

**PINNs**  
Pinns is a python library which creates neural networks that can solve differential equations.



**JAX-PI**

- Can also write your own code from scratch!

# Tutorial Exercise 1

- Estimate cerebral perfusion (=cerebral blood flow, CBF) in an infant using data from arterial spin labelling (ASL) MRI.
- Tutorial exercise GitHub page:



- [https://github.com/annien094/PINNs-tutorial-MICCAI-2024/blob/main/PINNs\\_ASL.ipynb](https://github.com/annien094/PINNs-tutorial-MICCAI-2024/blob/main/PINNs_ASL.ipynb)

# PINning Cerebral Blood Flow: Analysis of Perfusion MRI in Infants using Physics-Informed Neural Networks

**Christoforos Galazis**<sup>1,2,\*</sup>, **Ching-En Chiu**<sup>2,3</sup>, **Tomoki Arichi**<sup>4</sup>, **Anil A. Bharath**<sup>5</sup>  
**and Marta Varela**<sup>2</sup>

<sup>1</sup>*Department of Computing, Imperial College London, London, UK*

<sup>2</sup>*National Heart & Lung Institute, Imperial College London, London, UK*

<sup>3</sup>*Department of Electrical Engineering, Imperial College London, London, UK*

<sup>4</sup>*Centre for the Developing Brain, King's College London, London, UK*

<sup>5</sup>*Department of Bioengineering, Imperial College London, London, UK*

Correspondence\*:

Christoforos Galazis

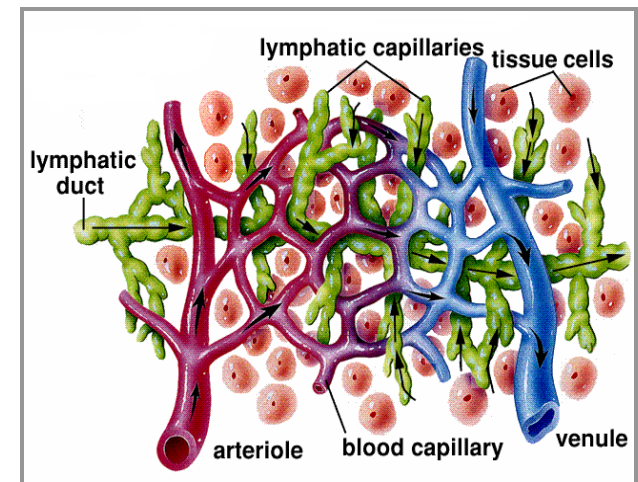
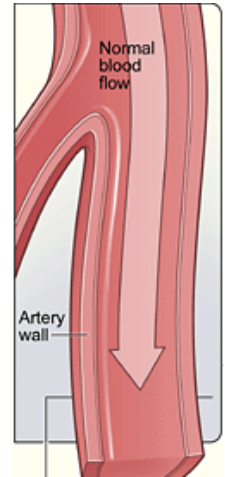
c.galazis20@imperial.ac.uk

(under review)

<https://github.com/cgalaz01/supinn>

# What is Perfusion?

- **Blood Flow in a Vessel:** volume of blood transported per unit time (mL blood/s)
- **Perfusion:** volume of blood delivered to a given amount of tissue per unit time (mL blood/100 g tissue/min)
- Information about:
  - Quality of local blood supply
  - Health/activity of the tissue

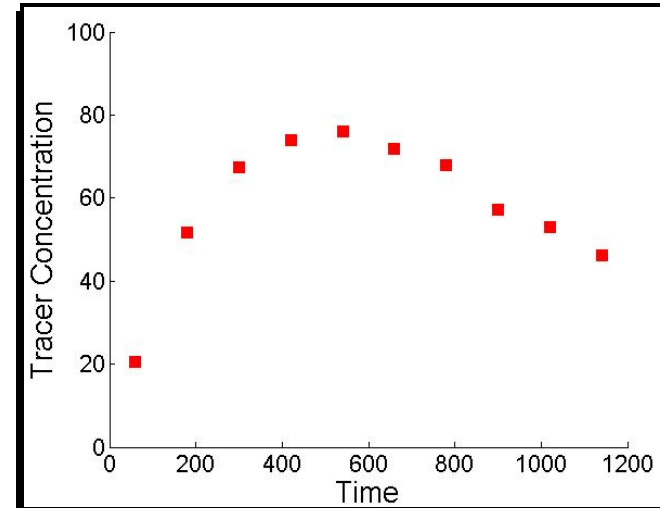


Cerebral Perfusion  $\Leftrightarrow$  Cerebral Blood Flow (CBF)



# How can Perfusion be Measured?

1. Inject blood tracer
2. Monitor concentration in tissue over time
3. Compute CBF using a model



# The ASL Experiment

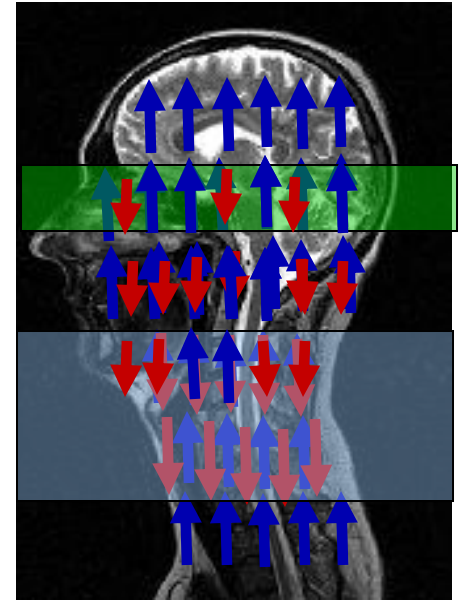
**1. Labelling:** Invert magnetization in the neck

- Labelled blood moves away
- Magnetization recovers with tissue-specific time constant  $T_1$

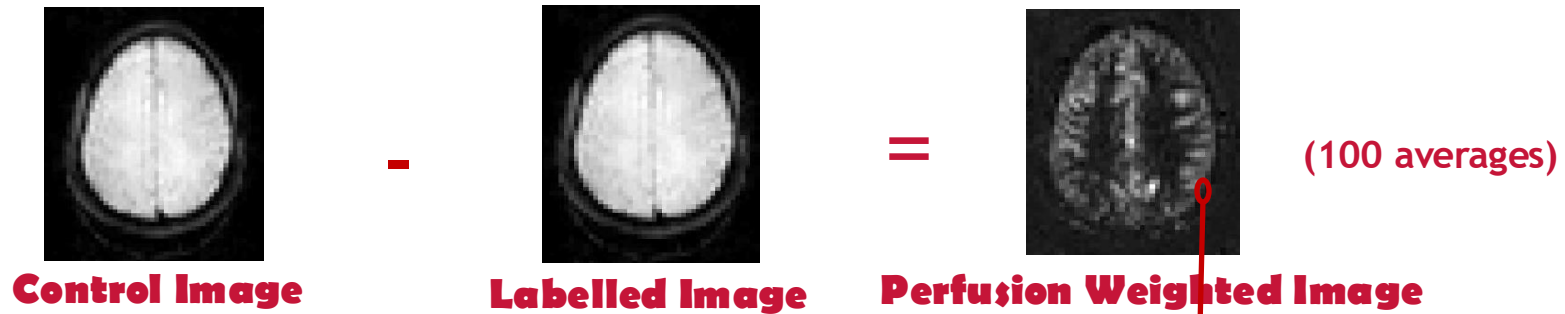
**2. Labelled Image:** Collect image in thin slice in the head

- Signal comes both from labelled blood and static tissue

**4. Control Image:** Acquire an image without labelling blood

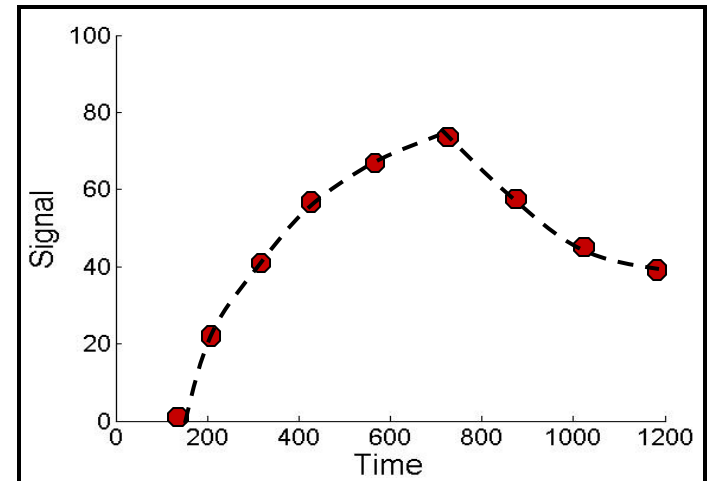


# The ASL Kinetic Model

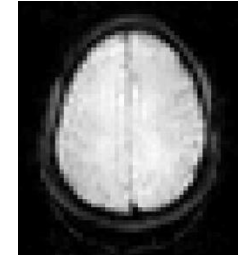
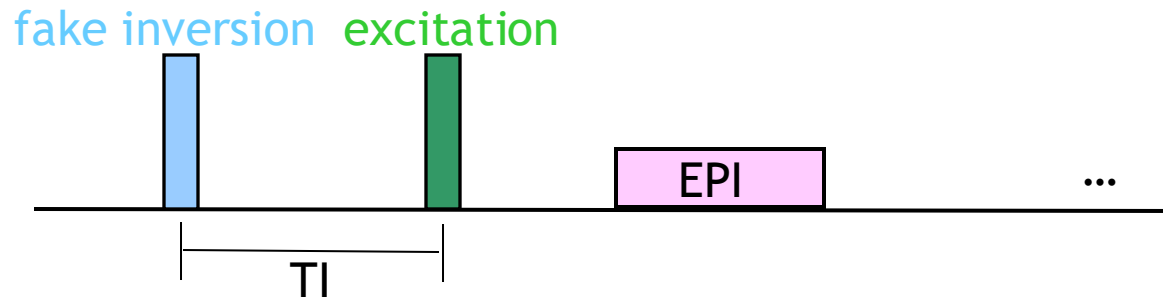


**5.** Repeat the experiment changing time gap between labelling and image acquisition

**6.** Fit data to kinetic model to extract CBF

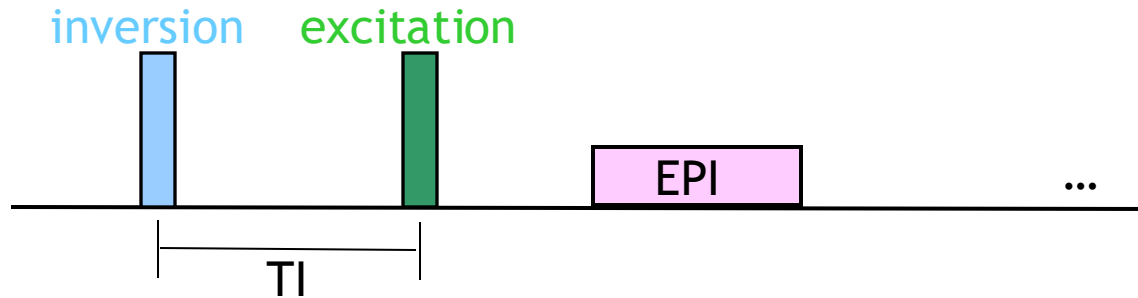


# Perfusion Weighted Images



**Control Image**

-



**Labelled Image**

=

5. Do many averages to increase SNR



**Perfusion Weighted Image**

(100 averages)

IMPERIAL

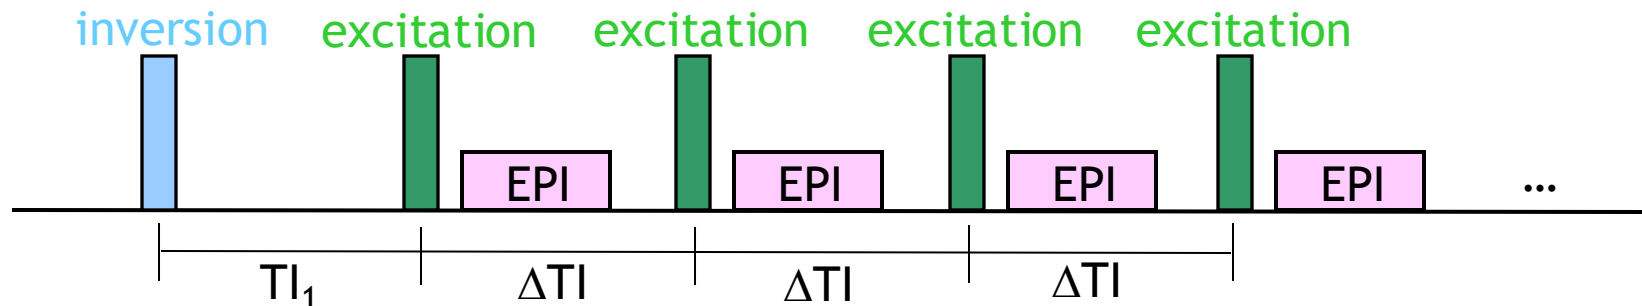
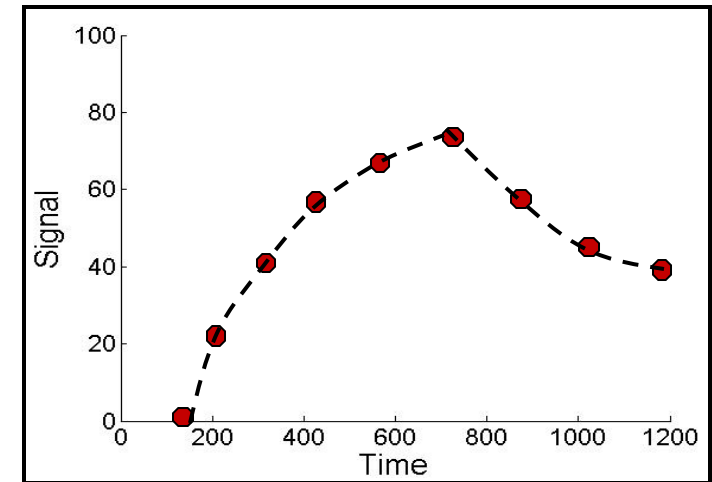
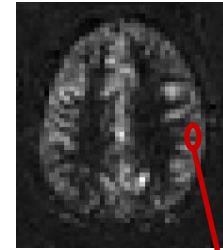
# Acquiring Images at Multiple TIs

5. Repeat the experiment changing time gap between labelling and image acquisition

6. Fit data to kinetic model to extract CBF

Or... faster (but giving lower SNR):

Perfusion Weighted Image



# (Simplified) ASL Model

Parameter		Healthy Adults	Neonates
$T_{1b}$	Longitudinal Recovery Time of Arterial Blood	1500 ms	Very variable
CBF	Cerebral Blood Flow	60 mL blood / 100 g tissue / min	Age dependent
$\tau$	Bolus Duration		
$\Delta t$	Bolus Arrival Time	-	

$$\frac{dS}{dt} = \begin{cases} 0 & \text{if } t < AT \\ CBF \times e^{-t/T_{1b}} \times \left(1 - \frac{t-AT}{T_{1b}}\right) & \text{if } AT \leq t < AT + \tau \\ -CBF \times e^{-t/T_{1b}} \times \frac{\tau}{T_{1b}} & \text{if } AT + \tau \leq t \end{cases}$$

# Tutorial Exercise 2

- Model the propagation of electrical signals in cardiac tissue in a 2D rectangular geometry.
- [https://github.com/annien094/PINNs-tutorial-MICCAI-2024/blob/main/PINNs\\_AP2D.ipynb](https://github.com/annien094/PINNs-tutorial-MICCAI-2024/blob/main/PINNs_AP2D.ipynb)



Scan for the  
exercise notebook

# IMPERIAL

## Physics-informed neural networks for cardiac electrophysiology in 3D and fibrillatory conditions

Annie Ching-En Chiu, Aditi Roy, Sarah Cechnicka,  
Arieh Levy, Ashvin Gupta, Christoforos Galazis, Kim Christensen,  
Danilo Mandic, Marta Varela

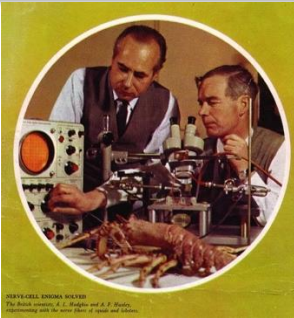
10<sup>th</sup> Oct 2024 @ STACOM workshop

<https://arxiv.org/pdf/2409.12712>

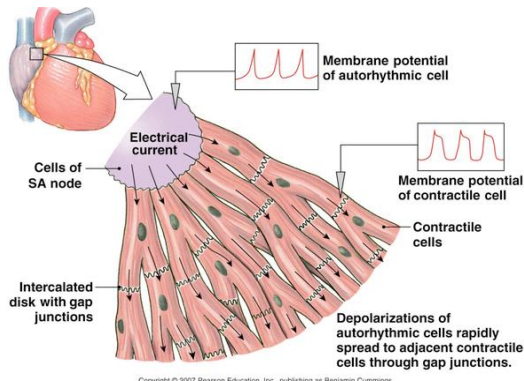
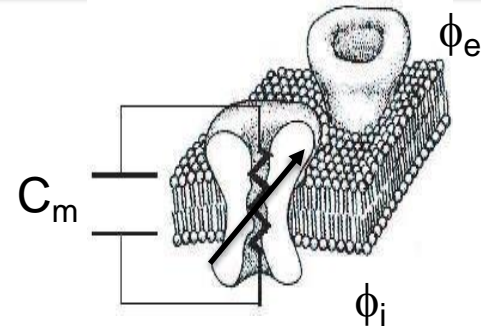




# Propagating Action Potentials



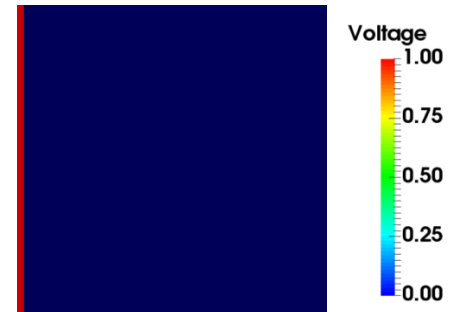
A. Hodgkin & A. Huxley,  
Nobel for Prize Medicine  
& Physiology, 1962



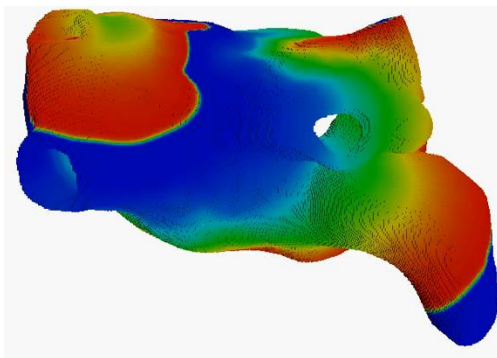
V “diffuses” to neighbouring cells.

$$\frac{\partial V}{\partial t} = \nabla \cdot (D \nabla V) - \frac{I_{ion}(V, t)}{C_m}$$

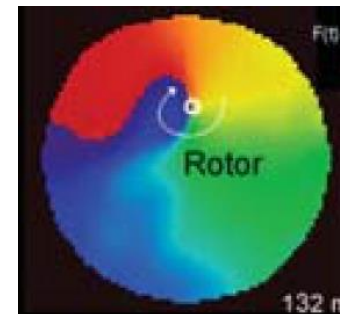
**D** : Diffusion tensor



**Travelling Wave Solutions**

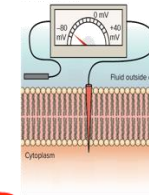
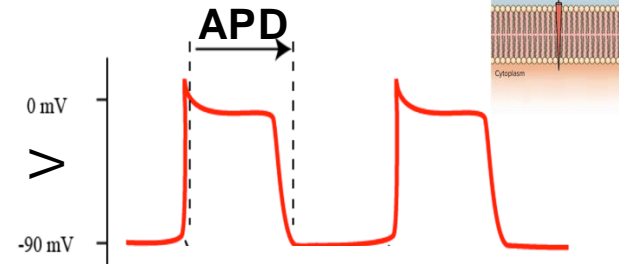
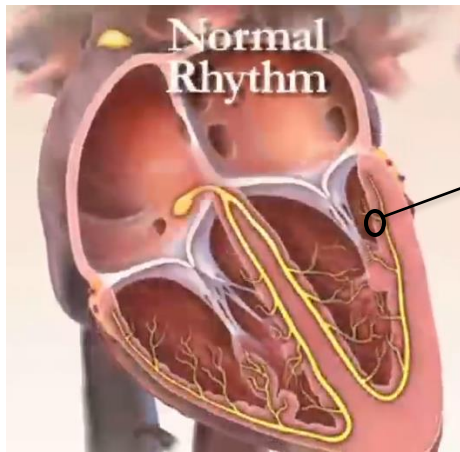


**Re-entrant waves  
(rotors) as drivers of  
arrhythmias.**



**IMPERIAL**

# Modelling Cardiac Electrophysiology



$$\frac{\partial V}{\partial t} = \vec{\nabla} \cdot (D \vec{\nabla} V) - kV(V - a)(V - 1) - VW$$

$$\frac{dW}{dt} = \left( \epsilon + \frac{\mu_1 W}{V + \mu_2} \right) (-W - kV(V - b - 1))$$

We can estimate:

- **D** is the (scalar) diffusion coefficient
- **a** “controls” the steepness of the action potential
- **b** “controls” action potential duration (APD)

